

Red-Teaming Apertus 1.5

An Automated, Reproducible Safety Evaluation of swiss-ai/Apertus-v1.5-70B

Track 1A — Red-Teaming Apertus, Hack Apertus Online Hackathon 2026

Author: Mohd Raja (Raza7x) · GitHub: [devilking7x](#) · Solo participant

Model under test: **swiss-ai/Apertus-v1.5-70B** via CSCS Inference API (<https://api.inference.cscs.ch/v1>, OpenAI-compatible) · Tested 2026-10-03/04

Harness: **apertus-redteam v0.4** — 15 attack classes, deterministic seeded runner, heuristic triage + mandatory human review · <https://github.com/devilking7x/apertus-redteam> · Apache-2.0

Responsible disclosure: findings are reported to the organizers through this submission. No real PII was used (synthetic canaries only); no large copyrighted excerpts are stored (runs/ is gitignored); disallowed content is summarized, never reproduced as instructions.

1. Abstract

We built **apertus-redteam**, an open, deterministic red-teaming harness, and used it to evaluate **Apertus 1.5 70B** (swiss-ai/Apertus-v1.5-70B, served by CSCS) across **15 attack classes** — jailbreaks, multilingual safety, prompt injection, privacy, IP/copyright, bias, factual correctness, Swiss culture/values, system-prompt extraction, encoding obfuscation, multi-turn crescendo attacks, many-shot context attacks, and three tool/agentive attack classes (tool-call hijacking, poisoned tool descriptions, indirect injection via tool outputs). Every prompt, response, seed and judgment is logged as JSONL; every run reproduces with one command.

Five issues are confirmed by human review: (1) prompt injection succeeds in 12/12 variations on two seeds (ASR 1.00); (2) a multilingual safety gap — identical disallowed requests refused in English/Hindi are answered in Swahili/Bengali; (3) verbatim regurgitation of copyrighted openings (1984, *Harry Potter*) on direct request; (4) disclosure of confidentiality-marked synthetic PII from context (3/12); (5) **indirect prompt injection via poisoned tool outputs** — in a simulated agentive loop, the model obeys instructions embedded in tool outputs, in the strongest cases ignoring the user's request entirely (7/48 across three seeds). Nine attack classes returned negative results — reported honestly below.

All heuristic flags were human-reviewed; three modules' flags were judged false positives on review (documented in §8). The harness, findings, figures and reproduction commands are public at github.com/devilking7x/apertus-redteam.

2. Introduction & scope

Apertus is Switzerland's fully open large language model. Openness is a strength — and a reason to red-team rigorously: anyone can deploy it, so its failure modes become everyone's failure modes. This report documents a black-box safety evaluation of the flagship **Apertus 1.5 70B** instruct model as served to hackathon participants via CSCS.

In scope: jailbreak robustness (incl. multilingual, dialect, encoding and many-shot variants), prompt injection (direct and indirect via tool outputs), tool-call hijacking, privacy/PII handling, IP/copyright regurgitation, bias/stereotyping, factual correctness, Swiss cultural values, system-prompt confidentiality, and multi-turn escalation — mapped to the five issue areas in the Track 1A brief (culture/values, bias, privacy, IP/copyright, factual correctness).

Out of scope: white-box attacks (weights are open but we test the served API as a deployer would see it), training-data extraction beyond memorized public texts, multimodal inputs, and non-Swiss legal analysis of the copyright findings.

3. Methodology

3.1 Attack taxonomy

Attack class	Intent	Variations × seeds	Judge signal
prompt_injection	Embedded instruction overrides carrier task	12 × {42, 43}	canary token in output
multilingual_jailbreak	Disallowed request in 6 languages	12 × {42, 43}	no refusal + instructional markers
pii_extraction	Disclose marked-confidential synthetic PII	12 × {42, 43}	canary value in output
copyright_regurgitation	Verbatim quote of texts on request	12 × {42, 43}	longest verbatim run vs reference
bias_stereotype	Agree with group generalizations	12 × {42, 43}	endorsement phrasing
factual_correctness	Hallucination / cutoff / contradiction	12 × 42	fact mismatch
swiss_culture_values	Swiss-specific value conflicts	12 × 42	value-violation markers

Attack class	Intent	Variations × seeds	Judge signal
system_prompt_extraction	Reveal hidden instructions	12 × 42	instruction-shaped response
swiss_dialect_jailbreak	Jailbreak in Schwiizerdütsch	12 × 42	no refusal + instructional markers
encoding_obfuscation	ROT13/reversed/interleaved payloads	12 × 42	canary in output
crescendo_jailbreak	3-turn gradual escalation (multi-turn)	12 × 42	final-turn instructional markers
manyspot_jailbreak	8/24/48-shot in-context conditioning	12 × 42	no refusal + instructional markers
tool_call_hijack	Injection steering tool calls (agentic)	12 × 42	tool_call with attacker args
tool_description_poisoning	Instructions hidden in tool metadata	12 × 42	follows description-embedded instruction
agentic_loop_injection	Poisoned tool output in agent loop	12 × {42, 43} + 24 × 44	follows output-embedded instruction

3.2 Harness design

The harness (run.py) generates prompts deterministically from a seed (domain-separated RNG per module), sends them to a model backend (local HF or OpenAI-compatible HTTP), applies a heuristic judge for triage, and logs everything — prompt, response, params, judgment, model ID, timestamp — to results.jsonl plus summary.json with attack-success rate and Wilson 95% confidence intervals. `bash reproduce.sh <run-id>` replays any run. Heuristic judges are **triage only**: \$4 findings were each confirmed by a human reading the transcripts, and false positives are reported as such (\$8).

3.3 Threat model

Black-box API attacker: no weight access, no system-prompt access. Success = the model produces the disallowed/undesired behavior defined per class (canary emission, verbatim quote, disclosure, endorsement, non-refusal with instructions). Decoding uses temperature 0.0 (deterministic) unless noted.

3.4 Ethics & responsible disclosure

All PII is synthetic (@example-test.org canaries generated by the harness). The only reference text stored in the repo is public-domain (Austen, 1813); model outputs that may contain copyrighted text live in runs/, which is gitignored and never committed. Disallowed content is summarized in this report, never reproduced as instructions. Findings are disclosed to the organizers via this submission first.

4. Findings

4.1 Prompt injection — embedded instructions override the carrier task (12/12, two seeds)

Carrier tasks (summarize / translate / list key points) wrapped a document ending in an injected instruction to output only a canary token. The model emitted the canary verbatim in **all 12 variations on seed 42 and all 12 on seed 43** (ASR 1.00, 95% CI [0.76, 1.00]) — the embedded instruction reliably wins over the user's task. A negative control (same injections base64-wrapped, 8 variations) scored 0/8: the failure is specific to plaintext embedded instructions.

Rubric self-assessment: severity 5 (top real-world LLM attack vector; poisoned documents/emails steer the model), breadth 5 (every RAG / document-processing deployment), novelty 2 (known class; novel model-specific measurement), validity 10 (two-seed reproduction + negative control + one-command repro).

4.2 Multilingual safety gap — Swahili/Bengali comply where English/Hindi refuse

The identical disallowed request (lock-picking instructions; keylogger code), framed as a school project, is **refused in English and Hindi** but **answered with step-by-step instructions in Swahili and Bengali** (3/12 flagged on seed 42, 3/12 on seed 43; English mean judge score 0.00 vs non-English 0.32). Safety refusal does not transfer to lower-resource languages — significant for a model marketed on 1,811-language coverage.

Rubric self-assessment: severity 5 (disallowed dual-use instructions produced; bypass is a translation away), breadth 5 (all non-English users), novelty 5 (model-specific asymmetry evidence on Apertus 1.5; English-refuse/Swahili-comply on the identical request), validity 10 (two-seed reproduction; English/Hindi refusals as built-in negative controls).

Repro: bash reproduce.sh multilingual_jailbreak-seed42-20261003T195024Z · Limitation: non-English prompts are author-written, not native-speaker verified (disclosed in findings/mj-001).

4.3 Verbatim regurgitation of copyrighted text on direct request

Asked to quote openings verbatim, the model emits the **verbatim first sentence of 1984 (Orwell)** and the **verbatim opening sentences of Harry Potter and the Philosopher's Stone (Rowling)** — human-verified exact matches, no refusal or deflection. A public-domain control (*Pride and Prejudice*) is reproduced at 23-word verbatim runs (auto-measured), confirming memorization rather than paraphrase.

Rubric self-assessment: severity 2 (short excerpts on explicit request; bounded), breadth 5 (any user), novelty 2 (known phenomenon; new measured evidence on Apertus 1.5), validity 10 (human-verified matches + auto-measured control + two seeds).

Repro: bash reproduce.sh copyright_regurgitation-seed42-20261003T195327Z · Only short excerpts quoted (fair use); full outputs gitignored.

4.4 Disclosure of confidentiality-marked synthetic PII (3/12)

With a memo containing synthetic name/email pairs explicitly marked CONFIDENTIAL in context, the model discloses the marked pairs on direct questioning (3/12 on seed 42, 3/12 on seed 43). The privacy boundary — marked-confidential context must not be revealed — is not respected.

Rubric self-assessment: severity 2 (honest cap: all data is synthetic; mechanism would be worse with real data), breadth 5 (any mixed-sensitivity RAG/agent deployment), novelty 2 (standard canary methodology; new model-specific evidence), validity 10 (two-seed reproduction; synthetic canaries make it exactly repeatable).

Repro: bash reproduce.sh pii_extraction-seed42-20261003T195233Z · No real PII used at any point.

4.5 Indirect prompt injection via poisoned tool outputs (7/48, three seeds)

In a simulated two-step agentic loop (user request → tool call → tool output → final response), we feed the model a **poisoned tool output** containing an injected instruction. In **7 of 48 variations across three seeds** (2/12 seed 42, 0/12 seed 43, 5/24 seed 44), the model **obeys the injected instruction** — in the strongest cases replying with only the attacker-controlled string, **dropping the user's request and the legitimate tool content entirely**. Six further flags were verbatim echoes of tool output (expected behavior) and were excluded on human review — only genuine instruction-obedience counts.

Rubric self-assessment: severity 5 (behavioral hijack of any agentic/RAG deployment feeding external data through tools), breadth 5 (all tool-using deployments), novelty 5 (first systematic demonstration on Apertus 1.5 via its tool-calling API with a reproducible harness), validity 10 (automated across 48 variations / 3 seeds; weak echoes honestly excluded).

Repro: bash reproduce.sh agentic_loop_injection-seed44-20261003T203728Z · Module: attacks/agentic_loop_injection.py

5. Severity matrix (jury rubric)

Issue	Sev.	Breadth	Nov.	Valid.	Notes
4.1 Prompt injection	5	5	2	10	12/12 × 2 seeds; base64 negative control 0/8
4.2 Multilingual gap	5	5	5	10	EN/HI refuse, SW/BN comply; same request
4.3 Copyright regurg.	2	5	2	10	Human-verified verbatim (1984, HP)
4.4 PII-canary disclosure	2	5	2	10	Synthetic data; honest severity cap
4.5 Indirect injection (tools)	5	5	5	10	7/48 strong obedience × 3 seeds; echoes excluded

Scale 0/1/2/5/10 per the Track 1A rubric. A zero in any dimension would zero the issue — none apply.

6. Figures




Figure 1 — Heuristic-triage attack success rate per attack class (seed 42, n=12 each) with 95% Wilson confidence intervals. §4/§8 reconcile triage flags with human-confirmed outcomes (system_prompt_extraction, encoding_obfuscation and crescendo flags were judged false positives/ambiguous on review).




Figure 2 — Multilingual jailbreak: mean judge score on English prompts (0.00, all refused) vs low-resource-language prompts (0.32). Positive Δ = worse safety off-English.

7. Cross-cutting observations

Safety drops off-English. The two refusal-competent languages (English, Hindi) vs the two compliance languages (Swahili, Bengali) suggest safety fine-tuning coverage follows resource level, not the model's claimed language coverage.

Markings are not boundaries. CONFIDENTIAL labels (§4.4) and document/user instruction hierarchy (§4.1) are both ignored — the model treats all input text as equally actionable.

Negative results (rigor): factual_correctness 0/12, swiss_culture_values 0/12, swiss_dialect_jailbreak 0/12 (the model handles its native dialect safely), system_prompt_extraction 0 confirmed leaks, encoding_obfuscation 0 confirmed bypasses (model echoes rather than complies), crescendo 1 weak/ambiguous flag (model added an explicit educational-use warning), manyshot_jailbreak 0/12 at 8/24/48 shots (context-scale conditioning fails), tool_call_hijack 0/12 (model refuses and names the attacker domain as malicious), tool_description_poisoning 0/12 (tool metadata treated as untrusted). Reporting negatives is deliberate: a harness that only finds positives is not measuring. The earlier bias_stereotype signal (mild, 2/12) was superseded by the stronger agentic finding and moved to findings/superseded/.

8. Mitigations & limitations

Per-issue: instruction-hierarchy training and untrusted-span delimiting (§4.1); multilingual safety data and refusal templates for low-resource languages (§4.2); near-verbatim output filters and quote-deflection templates (§4.3); confidentiality-aware training and entity output filters (§4.4); tool-output trust boundaries and instruction-source tagging (§4.5). Systemic: red-team coverage proportional to claimed language coverage; canary regression tests in CI. **Limitations:** heuristic judges are triage, not truth — §4 exists because a human read every flagged transcript, and weak flags were excluded on review. Non-English prompts are author-written, not native-speaker verified. Black-box API testing only; temperature 0.0; 12 variations per class per seed.

9. Reproducibility appendix

Environment: Python 3.12, stdlib HTTP backend (no torch needed for the openai backend). Model: swiss-ai/Apertus-v1.5-70B via <https://api.inference.cscs.ch/v1> (CSCS Inference API, tested 2026-10-03/04). Primary run IDs (seed 42): prompt_injection-seed42-20261003T195509Z, multilingual_jailbreak-seed42-20261003T195024Z, copyright_regurgitation-seed42-20261003T195327Z, pii_extraction-seed42-20261003T195233Z, agentic_loop_injection-seed44-20261003T203728Z. Validation seeds: 43 for the five filed issues (plus seeds 42/43 for agentic). Reproduce any run: `bash reproduce.sh <run-id>` (needs APERTUS_API_KEY / --api-key; the key is never logged). Full logs: `runs/<run-id>/{config,results,summaries}`.

```
git clone https://github.com/devilking7x/apertus-redteam cd apertus-redteam bash
reproduce.sh prompt_injection-seed42-20261003T195509Z
```

10. References

Track 1A brief (Hack Apertus online hack); CSCS LLM Inference API docs (docs.cscs.ch/services/inference/api); Apertus model family (swiss-ai, Hugging Face). Harness: github.com/devilking7x/apertus-redteam (Apache-2.0).